

CMPE 478: Parallel Processing
Spring 2019, Homework 1
 (due March 20th)
 (you can do this project in groups of at most 2 students)

In this project, you will use OpenMP to implement a parallel version of Google ranking process and apply it on the Erdos Web Graph which can be downloaded at <http://web-graph.org/>. The ranking will be done by carrying out the following iteration:

$$r^{(0)} = \begin{bmatrix} 1 \\ 1 \\ \dots \\ 1 \end{bmatrix}$$

Repeat
 $r^{(t+1)} = \alpha P r^{(t)} + (1 - \alpha) c$
until $\|r^{(t+1)} - r^{(t)}\|_1 \leq \varepsilon$

Here

- $c = \begin{bmatrix} 1 \\ 1 \\ \dots \\ 1 \end{bmatrix}$
- Take α as 0.2
- $\|r^{(t+1)} - r^{(t)}\|_1 = \sum |r_i^{(t+1)} - r_i^{(t)}|$
- ε is a small number, e.g. 10^{-6}

The matrix P is to be stored in CSR format. CSR format is explained below. You should provide a write-up of how you implemented your project and the following results:

- a) The timings obtained as shown below. **This table should be generated automatically by your program as a CSV file.**

Test No.	Scheduling Method	Chunk Size	No. of Iterations	Timings in secs for each number of threads							
				1	2	3	4	5	6	7	8
1											
2											
...											

- b) The names of the first 5 hosts that have the highest rankings.

Google Ranking Process

Details of Google ranking process is given in the following page:

- <http://infolab.stanford.edu/~backrub/google.html>

CSR Matrix Storage Format

Consider the following sparse matrix storage scheme, called compressed sparse row (CSR) format. An example of a matrix represented in this format is given below:

$$P = \begin{bmatrix} 11 & 0 & 13 & 14 & 0 \\ 0 & 0 & 23 & 24 & 0 \\ 31 & 32 & 33 & 34 & 0 \\ 0 & 42 & 0 & 44 & 0 \\ 51 & 52 & 0 & 0 & 55 \end{bmatrix}$$

The above matrix will be stored as follows:

```
row_begin = [ 0  3  5  9 11 14 ]
values    = [ 11 13 14 23 24 31 32 33 34 42 44 51 52 55 ]
col_indices = [ 0  2  3  2  3  0  1  2  3  1  3  0  1  4 ]
```

Let N stand for the number of nonzero entries in the matrix and n stand for the number of rows. The array **values** contains non-zero entries in the matrix in row wise order. The array **col_indices** gives the corresponding column indices of these values. The array **row_begin** of size $n+1$ stores the beginning index of each row in the **values** (and **col_indices** arrays). The last entry in **row_begin** stores $N+1$ so that the expression **row_begin**[$i+1$]-**row_begin**[i] gives the number of nonzeros in row i .

Grading

Your project will be graded according to the following criteria:

Documentation as a pdf file – a written document that includes : <ul style="list-style-type: none">• details of algorithm/ implementation,• details of the machine/CPU's you used ,• timings table (generated csv table)• discussion of results,.	12%
Comments in your code	8%
Implementation and tests	80%

Late Submission

If the project is submitted late, the following penalties will be applied:

- $0 < \text{hours late} \leq 24$: 25%
- $24 < \text{hours late} \leq 48$: 50%
- $\text{hours late} > 48$: 100%